

# Mannas Narang

[phone] | mannasnarang8@gmail.com | mannas.dev | github.com/mannasdev | x.com/mannasdev | Chandigarh, India

## ABOUT ME

---

I'm a self-taught engineer, still figuring most of it out as I go. I left a degree program at 18 and learned by freelancing, then spent about a year and a half helping build product at a US startup. These days I split my time between full-stack work and training models on a GPU server at home, usually a few steps behind the people I look up to but trying to close the gap. Based in Chandigarh, India.

## EXPERIENCE

---

### Wander

Remote

*Product Engineer II*

*Dec 2024 – Jul 2026*

- Designed and shipped the **WanderOS SEO Agent** to production: an autonomous agent that audits live host sites against Search Console and keyword data, proposes risk- and impact-tagged fixes, and ships approved changes directly to production sites – built around a five-level autonomy dial giving hosts graduated control from full manual review to full autonomy. Owned the agent's reliability post-launch, including failure debugging and guardrail design.
- Built and maintained third-party API integrations (Guesty, Hostaway, Lodgify, Track) that sync property data onto the marketplace in real time, including OAuth flows and event-driven sync pipelines handling data drift and partial failures across providers.
- Led a full-stack product redesign from first mock to production and shipped performance work that made key marketplace pages load near-instantly.

## AI & OPEN-SOURCE PROJECTS

---

### Vidura | *SQLite, sqlite-vec, macOS CLI*

2026

- Local, privacy-first tool that reads Claude Code session logs and surfaces recurring workflow-friction patterns via vector search over embedded session history – a working memory and retrieval layer built to evaluate agent behavior over time, not just log it.

### lerobot-lint | *Python, NumPy*

2026

- Open-source dataset quality linter for LeRobot: deterministic checks that catch dead joints, encoder wraparound, camera desync, and workspace-coverage gaps before they burn GPU hours – eslint for robot training data, and a hard-eval-gate mindset applied to a data pipeline.

### legal-bge-in | *PyTorch, Hugging Face*

2026

- Domain-specific embedding model fine-tuned on Indian case law for legal retrieval.  $\approx 25\%$  fewer missed cases and correct judgments ranked meaningfully higher vs. the base model (Recall@10 +7.2pts, MRR@10 +8.7pts on a held-out, leakage-checked eval set). Trained locally on a single consumer GPU (RTX 5060 Ti, 16GB), no paid APIs; released MIT-licensed.

## RESEARCH

---

### BOLI: A Dataset for Understanding Stuttering Experience and Analyzing Stuttered Speech

arXiv, 2025

- Co-authored a multilingual stuttered speech dataset built with and for people who stutter in India: read and spontaneous speech, five annotated stutter types, released open access.

## TECHNICAL SKILLS

---

**AI/ML:** Python, LLM APIs (prompting, structured outputs), PyTorch, Hugging Face (transformers, sentence-transformers), embedding fine-tuning, QLoRA, vector search / RAG (sqlite-vec), local inference (Ollama)  
**Backend & Systems:** TypeScript, Node.js, PostgreSQL, SQLite, REST APIs, OAuth & third-party SaaS integrations, event-driven architecture  
**Tooling & Deployment:** Git, GitHub Actions, Docker, Bash/Zsh, SSH, Linux GPU server administration